

A Comparative Study of Deep RL Algorithms: DQN, PPO, SAC, TD3

EPFL

Kevin Abou Jaoude · Youssef Dib · Fuad Khuri

EE-568 Reinforcement Learning · EPFL · Spring 2026

1. Four classic-control environments



2. Setup and taxonomy

We re-implement four canonical deep RL algorithms from scratch in PyTorch (no `stable-baselines3`) and benchmark them with 3 seeds each, 95% bootstrap CIs over seeds, and four pre-registered ablations.

	on-policy	off-policy
value-based	(none)	DQN
policy-based	PPO	SAC, TD3

Networks. (64, 64) MLPs for DQN, PPO; (256, 256) for SAC, TD3 (canonical paper widths). **Hardware.** RTX 3070 + EPFL gnoto V100.

3. Algorithms

DQN [3]

value-based, off-policy, discrete:

$$\mathcal{L} = \mathbb{E}[\ell_{\text{Hub}}(Q_{\theta} - r - \gamma(1-d)\max_{a'} Q_{\theta})]$$

ϵ -greedy, target sync every τ steps, Huber loss.

PPO [4]

policy-based, on-policy:

$$\mathcal{L} = -\mathbb{E}[\min(r_t \hat{A}_t, \text{clip}(r_t, 1 \pm \epsilon) \hat{A}_t)]$$

GAE advantages, clipped surrogate as a soft trust region.

SAC [2]

off-policy, max-entropy:

$$J = \mathbb{E}[\sum_t \gamma^t (r_t + \alpha \mathcal{H}[\pi(\cdot|s_t)])]$$

Squashed-Gaussian actor, auto α via $\log \alpha$, twin Q.

TD3 [1]

off-policy, deterministic:


$$y = r + \gamma(1-d) \min_i Q_{\phi_i}(s', \pi_{\hat{\theta}} + \hat{\epsilon})$$

Twin critics, target smoothing $\tilde{\sigma}=0.2$, delayed actor.


4. Implementation details that mattered

- DQN:** target-sync sweep on CartPole gives a **U-shape** ($\tau \in \{100, 1000, 5000\}$). The canonical $\tau=1000$ sits at the edge of stability, which explains the seed-bimodality.
- PPO:** at Pendulum truncation we fold $\gamma V(s_{t+1})$ into the reward so GAE accounts for the value of the truncated state; raw **done** flags bias returns by ~ 30 .
- SAC:** the tanh-Jacobian log-prob correction $-\sum_i \log(1 - \tanh^2 u_i)$ is mandatory; without it the entropy term is silently wrong and training never moves off the warmup floor.
- TD3:** the smoothed target action needs two clips, on the noise (± 0.5) and on the action range. Skipping the outer clip sticks the actor on the boundary.
- All four:** bootstrap mask uses **terminated** only, not **terminated or truncated**.


5. Cost, sample efficiency, model size



Key numbers at a glance




SAC sample-efficiency over PPO on Pendulum




Goal discoveries by SAC across α -sweep seeds on MCC

A. Learning curves: one figure per environment



CartPole. **DQN bimodal** (1/3 solve); PPO reliable, 380, CI [350, 402].

Acrobot. DQN solves cleanly in $\sim 30k$ steps; one PPO seed collapses.




Pendulum. **SAC and TD3:** $\sim 13\times$ more sample-efficient than PPO (SAC threshold at 7.4k vs PPO 95k).

MCC, sparse. **Only PPO ever finds the goal** (1/3 seeds, -2.5); SAC -0.8, TD3 -9.1.


B. Pre-registered ablations: turning the canonical knobs

Four pre-registered sweeps, one per algorithm, 3 seeds each.


Headline (MCC, sparse reward): **turning the off-policy exploration knobs (SAC α , TD3 σ) does not rescue them.**




SAC α sweep on MCC.
Auto, 0.05, 0.20: three failure modes, 0 goal discoveries / 9 seeds.



TD3 σ sweep on MCC.
 $\sigma \in \{0.1, 0.3, 0.5\}$: more noise is worse, not better.



DQN target-sync, CartPole.
 $\tau \in \{100, 1000, 5000\}$: **U-shape**, $\tau=1000$ at the edge of stability (explains seed-bimodality).



PPO clip ratio, Pendulum.
 $\epsilon \in \{0.1, 0.2, 0.3\}$: 0.1 too conservative; 0.2 & 0.3 tied.

Why off-policy continuous methods fail on MCC. Entropy regularization (SAC) and additive Gaussian noise (TD3) drive uniform coverage of the *local* action set, not directed exploration in state space. PPO's on-policy rollouts occasionally produce a long correlated action sequence that crosses the hill; off-policy replay samples are decorrelated and cannot.

C. Practical rule of thumb

- Discrete actions:** **DQN** wins on Acrobot, bimodal on CartPole; **PPO** is more reliable.
- Dense-reward continuous:** **SAC** and **TD3** dominate, 12.9 $\times$ over **PPO** on Pendulum.
- Sparse-reward continuous:** **PPO** is the only one to find the MCC goal; the failure is structural, not a tuning issue.

Pick SAC or TD3 first on continuous control; switch to PPO (or add explicit exploration bonuses) the moment reward becomes sparse.

References